

Chapter 4: Multiple Linear Regression

MATH 4230 Capstone Project

Dataset: Student Lifestyle, Mental Health and Burnout Insight

Adding all 19 predictors raises test R-squared to 0.707, with stress level, anxiety score, sleep hours, depression score, and social support emerging as the strongest standardized predictors of burnout severity.

Setup

```
In [1]: # — Standard Library —————
import os
import warnings
import itertools

# — Third-party —————
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
from scipy import stats
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import cross_val_score

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

# — Settings —————
warnings.filterwarnings('ignore')

RANDOM_STATE = 230
np.random.seed(RANDOM_STATE)

plt.rcParams['savefig.dpi'] = 150
plt.rcParams['axes.spines.top'] = False
plt.rcParams['axes.spines.right'] = False
sns.set_style('whitegrid')

# — Output directories —————
```

```
FIGURES_DIR = '../figures/ch04/'
RESULTS_DIR = '../results/ch04/'

os.makedirs(FIGURES_DIR, exist_ok=True)
os.makedirs(RESULTS_DIR, exist_ok=True)

print('Setup complete.')
print(f'RANDOM_STATE = {RANDOM_STATE}')
print(f'Figures -> {FIGURES_DIR}')
print(f'Results -> {RESULTS_DIR}')
```

```
Setup complete.
RANDOM_STATE = 230
Figures -> ../figures/ch04/
Results -> ../results/ch04/
```

§1. The Question

Which combination of academic behaviors, specifically study hours, exam pressure, and academic performance, best explains burnout severity when evaluated alongside demographic variables like gender and academic year?

§2. Method in Brief

Multiple linear regression extends simple linear regression by including several predictors at once. Each coefficient represents the partial effect of one variable on burnout while holding all other predictors constant, which lets us separate the individual contribution of each factor. The model is fit using ordinary least squares, and predictions can be made by plugging values into the resulting equation.

§3. Why This Method?

Chapter 3 showed that sleep hours alone explains only about 14% of the variation in burnout. Multiple linear regression lets us bring in the full set of available predictors and see how much of that remaining variation they account for. Each coefficient is still interpretable on its own, and we can directly compare which variables matter most after controlling for everything else. It also sets a strong linear baseline before moving into more complex methods in later chapters.

§4. Data Setup

```
In [2]: # — Load scaled arrays saved by Chapter 2 —————
# X_train and X_test are already StandardScaler-transformed (zero mean, unit variance)
# Scaling was fit on X_train only so no information from the test set leaked in.
# y values are the raw burnout_score (regression target), no scaling applied.

X_train = np.load('../results/X_train.npy')
X_test  = np.load('../results/X_test.npy')
y_train = np.load('../results/y_train_reg.npy')
y_test  = np.load('../results/y_test_reg.npy')

# — Load feature names —————
# feature_names.pkl holds the ordered list of column names matching the
# columns in X_train and X_test.
feature_names = joblib.load('../results/feature_names.pkl')

# — Load raw unscaled DataFrames for VIF and interpretation —————
# VIF is computed on unscaled data so the values are in their natural units.
# The scaled arrays are used for all model fitting and metrics.
df_train = pd.read_csv('../results/df_train.csv')
df_test  = pd.read_csv('../results/df_test.csv')

print(f'X_train shape      : {X_train.shape}')
print(f'X_test  shape      : {X_test.shape}')
print(f'y_train shape      : {y_train.shape}')
print(f'y_test  shape      : {y_test.shape}')
print(f'Number of features : {len(feature_names)}')
print()
print('Features:')
for i, name in enumerate(feature_names, 1):
    print(f' {i:2d}. {name}')
```

```
X_train shape      : (140000, 19)
X_test  shape      : (40000, 19)
y_train shape      : (140000,)
y_test  shape      : (40000,)
Number of features : 19
```

Features:

1. age
2. study_hours_per_day
3. exam_pressure
4. academic_performance
5. stress_level
6. anxiety_score
7. depression_score
8. sleep_hours
9. physical_activity
10. social_support
11. screen_time
12. internet_usage
13. financial_stress
14. family_expectation
15. gender_Male
16. gender_Other
17. academic_year_2
18. academic_year_3
19. academic_year_4

Table 4.1: Predictor justifications

The predictors below are the 19 standardized feature columns used in the multiple linear regression model.

§4a. Dummy Variable Encoding

```
In [3]: # — Identify dummy columns from feature_names —————
# Drop-first encoding was applied in Chapter 2 for both gender and academic_year.
# gender reference category      : Female
# academic_year reference       : Year 1 (Freshman)
# The dummy columns created are listed below.

dummy_cols = [f for f in feature_names if
               f.startswith('gender_') or f.startswith('academic_year_')]

print('Dummy columns created by drop-first encoding:')
for col in dummy_cols:
    print(f' {col}')

print()
print('Reference categories:')
print('  gender      -> Female (omitted category)')
print('  academic_year -> Year 1 / Freshman (omitted category)')
```

Dummy columns created by drop-first encoding:

gender_Male
gender_Other
academic_year_2
academic_year_3
academic_year_4

Reference categories:

gender -> Female (omitted category)
academic_year -> Year 1 / Freshman (omitted category)

```
In [4]: # — Fit a quick OLS on scaled data to pull dummy coefficients —————
# We fit the full model here so we can report the dummy coefficients.
# All values are reported to exactly 3 decimal places as required by Chapter 4.

X_train_df = pd.DataFrame(X_train, columns=feature_names)
X_sm       = sm.add_constant(X_train_df) # statsmodels needs an explicit intercept
ols_model  = sm.OLS(y_train, X_sm).fit()

# Extract dummy coefficients only
dummy_coef = pd.DataFrame({
    'Feature'      : dummy_cols,
    'Coefficient'  : [round(ols_model.params[col], 3) for col in dummy_cols],
    'Std Error'    : [round(ols_model.bse[col], 3) for col in dummy_cols],
    'p-value'     : [round(ols_model.pvalues[col], 3) for col in dummy_cols],
})

print('Table 4.2: Dummy variable coefficients relative to reference categories.')
print('All values to 3 decimal places.')
print()
print(dummy_coef.to_string(index=False))
```

Table 4.2: Dummy variable coefficients relative to reference categories.
All values to 3 decimal places.

Feature	Coefficient	Std Error	p-value
gender_Male	-0.003	0.002	0.160
gender_Other	-0.001	0.002	0.820
academic_year_2	-0.000	0.003	0.953
academic_year_3	-0.001	0.003	0.704
academic_year_4	-0.001	0.003	0.678

Interpretation of dummy coefficients.

Each dummy coefficient represents the difference in predicted burnout score for that group compared to the reference category, holding all other predictors constant. For example, `gender_Male` gives the average difference in burnout between male students and female students (the reference). A positive value means higher predicted burnout relative to the reference; a negative value means lower. The `academic_year` dummies compare each year group to first-year students.

§4b. Multicollinearity and VIF

```
In [5]: # — Compute VIF on unscaled data —————
# VIF (Variance Inflation Factor) measures how much the variance of a coefficient
# is inflated because of correlation with other predictors.
#
# Rule of thumb:
# VIF < 5: acceptable
# VIF 5-10: moderate concern
# VIF > 10: serious multicollinearity concern
#
# VIF is computed on unscaled data with dummy columns reconstructed from df_train.
# Scaling does not change VIF values, but using the raw DataFrame makes the
# variable names and units easier to reason about.

vif_df = df_train.drop(
    columns=['burnout_score', 'risk_level', 'mental_health_index', 'dropout_risk'],
    errors='ignore'
).copy()

# Apply the same dummy encoding as Chapter 2
vif_df = pd.get_dummies(vif_df, columns=['gender', 'academic_year'], drop_first=True)
vif_df = vif_df.astype(float)

# Reorder columns to match feature_names
vif_df = vif_df.reindex(columns=feature_names, fill_value=0)

# Compute VIF for each feature
vif_values = [
    variance_inflation_factor(vif_df.values, i)
    for i in range(vif_df.shape[1])
]

vif_table = pd.DataFrame({
    'Feature': feature_names,
    'VIF': [round(v, 3) for v in vif_values]
}).sort_values('VIF', ascending=False).reset_index(drop=True)

vif_table['Flag'] = vif_table['VIF'].apply(
    lambda v: 'HIGH (>10)' if v > 10 else ('MODERATE (>5)' if v > 5 else '')
)

# Main chapter table: only variables above the concern threshold
high_vif_table = vif_table[vif_table['VIF'] > 5].copy()

print('Table 4.3: Predictors with VIF greater than 5.')
display(high_vif_table)

# Full table saved for Appendix C
vif_full_csv = os.path.join(RESULTS_DIR, 'vif_full_table_appendix.csv')
vif_table.to_csv(vif_full_csv, index=False)

vif_high_csv = os.path.join(RESULTS_DIR, 'vif_high_predictors.csv')
```

```

high_vif_table.to_csv(vif_high_csv, index=False)

print()
print(f'Full VIF table saved for Appendix C: {vif_full_csv}')
print(f'Main high-VIF table saved: {vif_high_csv}')

```

Table 4.3: Predictors with VIF greater than 5.

	Feature	VIF	Flag
0	academic_performance	73.848	HIGH (>10)
1	exam_pressure	45.134	HIGH (>10)
2	stress_level	37.095	HIGH (>10)
3	screen_time	36.545	HIGH (>10)
4	age	34.052	HIGH (>10)
5	internet_usage	31.369	HIGH (>10)
6	sleep_hours	23.093	HIGH (>10)
7	study_hours_per_day	18.739	HIGH (>10)
8	family_expectation	13.532	HIGH (>10)
9	financial_stress	11.852	HIGH (>10)
10	anxiety_score	11.663	HIGH (>10)
11	social_support	9.294	MODERATE (>5)
12	physical_activity	5.596	MODERATE (>5)

Full VIF table saved for Appendix C: ../results/ch04/vif_full_table_appendix.csv
Main high-VIF table saved: ../results/ch04/vif_high_predictors.csv

```


In [6]: # — Figure 4.1: VIF bar chart —————
# Bars are colored red for VIF > 10, orange for VIF 5-10, steelblue otherwise.

colors = [
    'firebrick' if v > 10 else
    'darkorange' if v > 5 else
    'steelblue'
    for v in vif_table['VIF']
]

fig, ax = plt.subplots(figsize=(9, 6))
ax.barh(vif_table['Feature'][::-1], vif_table['VIF'][::-1], color=colors[::-1])
ax.axvline(5, color='darkorange', linestyle='--', linewidth=1.2, label='VIF = 5')
ax.axvline(10, color='firebrick', linestyle='--', linewidth=1.2, label='VIF = 10')
ax.set_xlabel('Variance Inflation Factor (VIF)', fontsize=11)
ax.set_title('Figure 4.1: VIF for All Predictors', fontsize=11, pad=10)
ax.legend(fontsize=9)
plt.tight_layout()
fig_path_41 = os.path.join(FIGURES_DIR, 'figure_4_1_vif.png')
plt.savefig(fig_path_41, dpi=150, bbox_inches='tight')

```

```
plt.show()
print(f'Saved: {fig_path_41}')
```



Saved: ../figures/ch04/figure_4_1_vif.png

§4c. Subset Selection and Information Criteria

```
In [7]: # — Forward stepwise selection using AIC —————
# Forward stepwise selection starts with no predictors and adds one variable
# at a time, choosing the one that lowers AIC the most at each step.
# AIC penalizes model complexity, so it stops rewarding additional predictors
# once the improvement in fit no longer offsets the added complexity.
#
# statsmodels OLS is used here because it natively reports AIC and BIC.
# Mallows Cp is computed manually:
#   Cp = (RSS_p / sigma^2) - (n - 2p)
# where sigma^2 is estimated from the full model.

n          = X_train.shape[0]
p_full     = X_train.shape[1]
X_train_df = pd.DataFrame(X_train, columns=feature_names)

# Fit full model first to get sigma^2 for Mallows Cp
X_full_sm  = sm.add_constant(X_train_df)
full_model = sm.OLS(y_train, X_full_sm).fit()
sigma2     = full_model.mse_resid # MSE of full model as sigma^2 estimate

remaining  = list(feature_names)
selected   = []
results    = []
```

```

for step in range(1, p_full + 1):
    best_aic = np.inf
    best_var = None
    best_fit = None

    for candidate in remaining:
        cols = selected + [candidate]
        X_sub = sm.add_constant(X_train_df[cols])
        fit = sm.OLS(y_train, X_sub).fit()
        if fit.aic < best_aic:
            best_aic = fit.aic
            best_var = candidate
            best_fit = fit

    selected.append(best_var)
    remaining.remove(best_var)

# Mallows Cp
rss_p = best_fit.ssr
cp = round((rss_p / sigma2) - (n - 2 * step), 3)

results.append({
    'Step' : step,
    'Variable Added' : best_var,
    'Num Predictors' : step,
    'AIC' : round(best_fit.aic, 3),
    'BIC' : round(best_fit.bic, 3),
    'Mallows Cp' : cp,
})

# Stop early if AIC starts increasing
if len(results) > 1 and results[-1]['AIC'] > results[-2]['AIC']:
    print(f'AIC increased at step {step}. Stopping early.')
    break

subset_df = pd.DataFrame(results)
print('Forward Stepwise Selection Results (all values to 3 decimal places):')
print(subset_df.to_string(index=False))

```

AIC increased at step 11. Stopping early.

Forward Stepwise Selection Results (all values to 3 decimal places):

Step	Variable Added	Num Predictors	AIC	BIC	Mallows Cp
1	stress_level	1	422209.705	422229.404	66673.930
2	social_support	2	404092.758	404122.307	41586.775
3	sleep_hours	3	391570.277	391609.675	26049.704
4	anxiety_score	4	382475.815	382525.062	15606.085
5	depression_score	5	375107.780	375166.876	7628.788
6	exam_pressure	6	367693.773	367762.719	14.701
7	screen_time	7	367691.257	367770.052	12.185
8	internet_usage	8	367687.300	367775.944	8.228
9	family_expectation	9	367684.880	367783.374	5.809
10	study_hours_per_day	10	367683.212	367791.555	4.140
11	gender_Male	11	367683.293	367801.486	4.222

```
In [8]: # — Save subset selection table —————
subset_csv = os.path.join(RESULTS_DIR, 'subset_selection.csv')
subset_df.to_csv(subset_csv, index=False)
print(f'Saved: {subset_csv}')

# — Identify selected model —————
best_row = subset_df.loc[subset_df['AIC'].idxmin()]
n_selected = int(best_row['Num Predictors'])
print(f'\nSelected model: {n_selected} predictors (lowest AIC = {best_row["AIC"]:.3
```

Saved: ../results/ch04/subset_selection.csv

Selected model: 10 predictors (lowest AIC = 367683.212)

```
In [9]: # — Figure 4.2: AIC and BIC vs number of predictors —————
# AIC and BIC both penalize model complexity, but BIC penalizes complexity more heavily
# This figure marks the model size selected by each criterion.

aic_best_row = subset_df.loc[subset_df['AIC'].idxmin()]
bic_best_row = subset_df.loc[subset_df['BIC'].idxmin()]

n_aic_best = int(aic_best_row['Num Predictors'])
n_bic_best = int(bic_best_row['Num Predictors'])

fig, ax = plt.subplots(figsize=(9, 5))

ax.plot(
    subset_df['Num Predictors'],
    subset_df['AIC'],
    marker='o',
    color='steelblue',
    label='AIC'
)

ax.plot(
    subset_df['Num Predictors'],
    subset_df['BIC'],
    marker='s',
    color='firebrick',
    label='BIC'
)

ax.axvline(
    n_aic_best,
    color='steelblue',
    linestyle='--',
    linewidth=1,
    alpha=0.65,
    label=f'Min AIC at k={n_aic_best}'
)

ax.axvline(
    n_bic_best,
    color='firebrick',
    linestyle=':',
    linewidth=1.3,
```

```

alpha=0.75,
label=f'Min BIC at k={n_bic_best}'
)

ax.set_xlabel('Number of Predictors', fontsize=11)
ax.set_ylabel('Information Criterion', fontsize=11)
ax.set_title(
    'Figure 4.2: AIC and BIC vs Number of Predictors',
    fontsize=11,
    pad=10
)
ax.legend(fontsize=9)


plt.tight_layout()

fig_path_42 = os.path.join(FIGURES_DIR, 'figure_4_2_aic_bic_selection.png')
plt.savefig(fig_path_42, dpi=150, bbox_inches='tight')
plt.show()

print(f'Saved: {fig_path_42}')
print(f'Minimum AIC model size: {n_aic_best} predictors')
print(f'Minimum BIC model size: {n_bic_best} predictors')

```

Figure 4.2: AIC and BIC vs Number of Predictors



Saved: ../figures/ch04/figure_4_2_aic_bic_selection.png
 Minimum AIC model size: 10 predictors
 Minimum BIC model size: 6 predictors

Selected model justification.

The forward stepwise procedure selected the model at the step with the lowest AIC. AIC balances goodness of fit against the number of parameters, so it naturally penalizes adding predictors that do not contribute meaningfully. BIC selects a smaller 6-predictor model because it penalizes complexity more heavily. The full 19-predictor model is still reported in §5 because it gives the clearest coefficient table for all variables and serves as the main multiple linear regression baseline.

Greedy Best-Subset Approximation

Because there are 19 predictors, exhaustive best subset selection would require evaluating 524,288 possible models. To keep the analysis computationally practical, this section uses a greedy best-subset approximation based on the forward-selection order. This does not test every possible subset, so it should be interpreted as a model-selection comparison rather than a true exhaustive best-subset search.

```
In [10]: # — Greedy best-subset approximation —————
# This is not true exhaustive best subset selection.
# It uses the forward-selected variable order as a practical approximation.
#
# With 19 predictors, true exhaustive best subset selection would require
# evaluating  $2^{19} = 524,288$  possible models. This approximation evaluates
# the nested models created by the forward-selected ordering.

X_train_df = pd.DataFrame(X_train, columns=feature_names)

best_subset_results = []
ordered_vars = subset_df['Variable Added'].tolist()

for k in range(1, len(ordered_vars) + 1):
    cols = ordered_vars[:k]

    X_sub = sm.add_constant(X_train_df[cols])
    fit = sm.OLS(y_train, X_sub).fit()

    rss = fit.ssr
    cp = round((rss / sigma2) - (n - 2 * k), 3)

    best_subset_results.append({
        'Num Predictors': k,
        'Variables': cols,
        'BIC': round(fit.bic, 3),
        'AIC': round(fit.aic, 3),
        'Mallows Cp': cp,
    })

best_subset_df = pd.DataFrame(best_subset_results)

# Save full table for Appendix C before printing the file path
best_subset_csv = os.path.join(
    RESULTS_DIR,
    'greedy_best_subset_table_appendix.csv'
)

best_subset_df.to_csv(best_subset_csv, index=False)

# Select model with lowest BIC
best_bs_row = best_subset_df.loc[best_subset_df['BIC'].idxmin()]
n_best_subset = int(best_bs_row['Num Predictors'])
vars_best_subset = best_bs_row['Variables']
```

```

print('Greedy best-subset approximation using BIC:')
print(best_subset_df[['Num Predictors', 'BIC', 'AIC', 'Mallows Cp']].to_string(index=False))

print(f'\nSelected model size : {n_best_subset} predictors (lowest BIC = {best_bs_r})')
print(f'Full greedy best-subset table saved for Appendix C: {best_subset_csv}')
print(f'Selected predictors : {list(vars_best_subset)}')

# Fit selected model and get coefficients
X_bs_sm = sm.add_constant(X_train_df[vars_best_subset])
bs_fit = sm.OLS(y_train, X_bs_sm).fit()

print('\nCoefficients:')
for feat in vars_best_subset:
    print(f' {feat:<25s}: {bs_fit.params[feat]:+.5f}')

# Test MSE
X_test_df = pd.DataFrame(X_test, columns=feature_names)
X_bs_test = sm.add_constant(X_test_df[vars_best_subset])

y_pred_bs = bs_fit.predict(X_bs_test)

bs_test_mse = round(float(np.mean((y_test - y_pred_bs) ** 2)), 5)
bs_test_rmse = round(float(np.sqrt(bs_test_mse)), 5)

print(f'\nTest MSE : {bs_test_mse}')
print(f'Test RMSE : {bs_test_rmse}')

```

Greedy best-subset approximation using BIC:

Num Predictors	BIC	AIC	Mallows Cp
1	422229.404	422209.705	66673.930
2	404122.307	404092.758	41586.775
3	391609.675	391570.277	26049.704
4	382525.062	382475.815	15606.085
5	375166.876	375107.780	7628.788
6	367762.719	367693.773	14.701
7	367770.052	367691.257	12.185
8	367775.944	367687.300	8.228
9	367783.374	367684.880	5.809
10	367791.555	367683.212	4.140
11	367801.486	367683.293	4.222

Selected model size : 6 predictors (lowest BIC = 367762.719)

Full greedy best-subset table saved for Appendix C: ../results/ch04/greedy_best_subset_table_appendix.csv

Selected predictors : ['stress_level', 'social_support', 'sleep_hours', 'anxiety_score', 'depression_score', 'exam_pressure']

Coefficients:

stress_level	: +0.57687
social_support	: -0.26726
sleep_hours	: -0.33928
anxiety_score	: +0.37617
depression_score	: +0.30996
exam_pressure	: +0.23852

Test MSE : 0.81065

Test RMSE : 0.90036

```
In [11]: # — Appendix Figure C.4.1: Greedy best-subset BIC vs model size —————
# This figure supports the model-selection discussion but is placed in Appendix C
# because Figure 4.2 already gives the main AIC/BIC selection result.

fig, ax = plt.subplots(figsize=(8, 5))

ax.plot(
    best_subset_df['Num Predictors'],
    best_subset_df['BIC'],
    'o-',
    color='steelblue',
    linewidth=2,
    label='BIC'
)

ax.plot(
    n_best_subset,
    best_bs_row['BIC'],
    'o',
    color='firebrick',
    markersize=10,
    label=f'Min BIC at k={n_best_subset}'
)

ax.set_xlabel('Number of Predictors', fontsize=11)
```

```


ax.set_ylabel('BIC', fontsize=11)
ax.set_title(
    'Appendix Figure C.4.1: Greedy Best-Subset BIC vs Model Size',
    fontsize=11,
    pad=10
)
ax.legend(fontsize=9)

plt.tight_layout()

fig_app_c41 = os.path.join(FIGURES_DIR, 'appendix_c_4_1_greedy_best_subset_bic.png')
plt.savefig(fig_app_c41, dpi=150, bbox_inches='tight')
plt.show()

print(f'Saved: {fig_app_c41}')

```



Saved: ../figures/ch04/appendix_c_4_1_greedy_best_subset_bic.png

Backward Stepwise Selection

```

In [12]: # — Backward stepwise selection using BIC —————
# Backward stepwise starts with ALL predictors included and removes one
# at a time, choosing the removal that lowers BIC the most at each step.
# It stops when removing any variable would increase BIC.
# This complements forward stepwise: forward builds up from nothing,
# backward prunes down from everything.

remaining_bwd = list(feature_names)
selected_bwd = list(feature_names)
results_bwd = []

# Record full model first
X_full_bwd = sm.add_constant(X_train_df[selected_bwd])

```

```

fit_full = sm.OLS(y_train, X_full_bwd).fit()

results_bwd.append({
    'Step': 0,
    'Variable Removed': 'None (full model)',
    'Num Predictors': len(selected_bwd),
    'BIC': round(fit_full.bic, 3),
    'AIC': round(fit_full.aic, 3),
})

for step in range(1, len(feature_names)):
    best_bic_bwd = np.inf
    worst_var = None
    best_fit_bwd = None

    for candidate in selected_bwd:
        cols = [v for v in selected_bwd if v != candidate]
        X_sub = sm.add_constant(X_train_df[cols])
        fit = sm.OLS(y_train, X_sub).fit()

        if fit.bic < best_bic_bwd:
            best_bic_bwd = fit.bic
            worst_var = candidate
            best_fit_bwd = fit

    # Stop if removing any variable increases BIC
    if best_bic_bwd >= results_bwd[-1]['BIC']:
        print(f'Backward stepwise stopped at step {step}: removing any variable inc
        break

    selected_bwd.remove(worst_var)

    results_bwd.append({
        'Step': step,
        'Variable Removed': worst_var,
        'Num Predictors': len(selected_bwd),
        'BIC': round(best_bic_bwd, 3),
        'AIC': round(best_fit_bwd.aic, 3),
    })

bwd_df = pd.DataFrame(results_bwd)

# Save full table for Appendix C before printing the file path
bwd_csv = os.path.join(
    RESULTS_DIR,
    'backward_stepwise_table_appendix.csv'
)

bwd_df.to_csv(bwd_csv, index=False)

best_bwd_row = bwd_df.loc[bwd_df['BIC'].idxmin()]
n_best_bwd = int(best_bwd_row['Num Predictors'])
vars_best_bwd = selected_bwd

print('Backward Stepwise Results:')
print(bwd_df[['Step', 'Variable Removed', 'Num Predictors', 'BIC', 'AIC']].to_strin

```

```

print(f'\nSelected model size : {n_best_bwd} predictors (lowest BIC = {best_bwd_row}
print(f'Selected predictors : {vars_best_bwd}')
print(f'Full backward stepwise table saved for Appendix C: {bwd_csv}')

# Fit selected model
X_bwd_sm = sm.add_constant(X_train_df[vars_best_bwd])
bwd_fit = sm.OLS(y_train, X_bwd_sm).fit()

print('\nCoefficients:')
for feat in vars_best_bwd:
    print(f' {feat:<25s}: {bwd_fit.params[feat]:+.5f}')

# Test MSE
X_bwd_test = sm.add_constant(X_test_df[vars_best_bwd])
y_pred_bwd = bwd_fit.predict(X_bwd_test)

bwd_test_mse = round(float(np.mean((y_test - y_pred_bwd) ** 2)), 5)
bwd_test_rmse = round(float(np.sqrt(bwd_test_mse)), 5)

print(f'\nTest MSE : {bwd_test_mse}')
print(f'Test RMSE : {bwd_test_rmse}')

```

Backward stepwise stopped at step 14: removing any variable increases BIC.

Backward Stepwise Results:

Step	Variable Removed	Num Predictors	BIC	AIC
0	None (full model)	19	367894.059	367697.071
1	academic_year_2	18	367882.213	367695.074
2	gender_Other	17	367870.415	367693.126
3	academic_performance	16	367858.632	367691.192
4	age	15	367846.856	367689.266
5	academic_year_3	14	367835.173	367687.432
6	academic_year_4	13	367823.432	367685.540
7	financial_stress	12	367812.429	367684.387
8	physical_activity	11	367801.486	367683.293
9	gender_Male	10	367791.555	367683.212
10	study_hours_per_day	9	367783.374	367684.880
11	family_expectation	8	367775.944	367687.300
12	internet_usage	7	367770.052	367691.257
13	screen_time	6	367762.719	367693.773

Selected model size : 6 predictors (lowest BIC = 367762.719)

Selected predictors : ['exam_pressure', 'stress_level', 'anxiety_score', 'depression_score', 'sleep_hours', 'social_support']

Full backward stepwise table saved for Appendix C: ../results/ch04/backward_stepwise_table_appendix.csv

Coefficients:

exam_pressure	: +0.23852
stress_level	: +0.57687
anxiety_score	: +0.37617
depression_score	: +0.30996
sleep_hours	: -0.33928
social_support	: -0.26726

Test MSE : 0.81065

Test RMSE : 0.90036

```
In [13]: # — Appendix Figure C.4.2: Backward stepwise BIC vs model size —————

fig, ax = plt.subplots(figsize=(8, 5))

ax.plot(
    bwd_df['Num Predictors'],
    bwd_df['BIC'],
    'o-',
    color='darkorange',
    linewidth=2,
    label='BIC'
)

ax.plot(
    n_best_bwd,
    best_bwd_row['BIC'],
    'o',
    color='firebrick',
    markersize=10,
    label=f'Min BIC at k={n_best_bwd}'
)
```

```


ax.set_xlabel('Number of Predictors', fontsize=11)
ax.set_ylabel('BIC', fontsize=11)
ax.set_title(
    'Appendix Figure C.4.2: Backward Stepwise BIC vs Model Size',
    fontsize=11,
    pad=10
)
ax.legend(fontsize=9)

plt.tight_layout()

fig_app_c42 = os.path.join(FIGURES_DIR, 'appendix_c_4_2_backward_stepwise_bic.png')
plt.savefig(fig_app_c42, dpi=150, bbox_inches='tight')
plt.show()

print(f'Saved: {fig_app_c42}')

```



Saved: ../figures/ch04/appendix_c_4_2_backward_stepwise_bic.png

Appendix C support plot: Forward Stepwise BIC

This plot is saved for Appendix C because Figure 4.2 already summarizes the main AIC and BIC selection behavior.

```

In [14]: # — Appendix Figure C.4.3: Forward stepwise BIC vs model size —————
# This reuses subset_df from the forward stepwise procedure.

fwd_best_row = subset_df.loc[subset_df['BIC'].idxmin()]
n_fwd_best = int(fwd_best_row['Num Predictors'])

fig, ax = plt.subplots(figsize=(8, 5))

```

```

ax.plot(
    subset_df['Num Predictors'],
    subset_df['BIC'],
    'o-',
    color='steelblue',
    linewidth=2,
    label='BIC'
)

ax.plot(
    n_fwd_best,
    fwd_best_row['BIC'],
    'o',
    color='firebrick',
    markersize=10,
    label=f'Min BIC at k={n_fwd_best}'
)

ax.set_xlabel('Number of Predictors', fontsize=11)
ax.set_ylabel('BIC', fontsize=11)
ax.set_title(
    'Appendix Figure C.4.3: Forward Stepwise BIC vs Model Size',
    fontsize=11,
    pad=10
)
ax.legend(fontsize=9)


plt.tight_layout()

fig_app_c43 = os.path.join(FIGURES_DIR, 'appendix_c_4_3_forward_stepwise_bic.png')
plt.savefig(fig_app_c43, dpi=150, bbox_inches='tight')
plt.show()

print(f'Saved: {fig_app_c43}')

```

Appendix Figure C.4.3: Forward Stepwise BIC vs Model Size



Saved: ../figures/ch04/appendix_c_4_3_forward_stepwise_bic.png

Three-Method Comparison Summary

```
In [15]: # — Forward stepwise test MSE (for comparison) —————
# Recompute using the forward-selected predictors from subset_df
vars_fwd      = subset_df.loc[:,subset_df['AIC'].idxmin(), 'Variable Added'].tolist()
X_fwd_sm      = sm.add_constant(X_train_df[vars_fwd])
fwd_fit_final = sm.OLS(y_train, X_fwd_sm).fit()
X_fwd_test     = sm.add_constant(X_test_df[vars_fwd])
y_pred_fwd     = fwd_fit_final.predict(X_fwd_test)
fwd_test_mse   = round(float(np.mean((y_test - y_pred_fwd)**2)), 5)
fwd_test_rmse  = round(float(np.sqrt(fwd_test_mse)), 5)

# — Summary comparison table —————
def formula(target, preds):
    return f'burnout ~ {" + ".join(preds)}'

summary_data = {
    'Method'                : ['Forward Stepwise', 'Best Subset (greedy)', ''],
    'Predictors Selected'    : [len(vars_fwd), n_best_subset, n_best_bwd],
    'BIC at Selection'       : [round(fwd_best_row['BIC'], 3),
                                round(best_bs_row['BIC'], 3),
                                round(best_bwd_row['BIC'], 3)],
    'Test MSE'               : [fwd_test_mse, bs_test_mse, bwd_test_mse],
    'Test RMSE'              : [fwd_test_rmse, bs_test_rmse, bwd_test_rmse],
    'Formula'                : [formula('burnout', vars_fwd),
                                formula('burnout', list(vars_best_subset)),
                                formula('burnout', vars_best_bwd)],
}

summary_df = pd.DataFrame(summary_data)
print('Table 4.4: Subset selection comparison summary.')
```

```

print('=' * 100)
print(summary_df[['Method', 'Predictors Selected', 'BIC at Selection', 'Test MSE', 'Test RMSE']])
print('=' * 100)
print()
print('Selected formulas:')
for _, row in summary_df.iterrows():
    print(f' {row["Method"]}: {row["Formula"]}')

# Save
comparison_csv = os.path.join(RESULTS_DIR, 'table_4_4_subset_comparison.csv')
summary_df.to_csv(comparison_csv, index=False)
print(f'\nSaved: {comparison_csv}')

# Agreement check
sets = [set(vars_fwd), set(vars_best_subset), set(vars_best_bwd)]
all_agree = sets[0] == sets[1] == sets[2]
print(f'\nAll three methods selected identical predictors: {all_agree}')
best_mse_method = summary_df.loc[summary_df['Test MSE'].idxmin(), 'Method']
print(f'Lowest test MSE: {best_mse_method} ({summary_df["Test MSE"].min():.5f})')

```

Table 4.4: Subset selection comparison summary.

	Method	Predictors Selected	BIC at Selection	Test MSE	Test RMSE
	Forward Stepwise	10	367762.719	0.81063	0.90035
Best Subset (greedy)		6	367762.719	0.81065	0.90036
	Backward Stepwise	6	367762.719	0.81065	0.90036

Selected formulas:

Forward Stepwise: burnout ~ stress_level + social_support + sleep_hours + anxiety_score + depression_score + exam_pressure + screen_time + internet_usage + family_expectation + study_hours_per_day

Best Subset (greedy): burnout ~ stress_level + social_support + sleep_hours + anxiety_score + depression_score + exam_pressure

Backward Stepwise: burnout ~ exam_pressure + stress_level + anxiety_score + depression_score + sleep_hours + social_support

Saved: ../results/ch04/table_4_4_subset_comparison.csv

All three methods selected identical predictors: False

Lowest test MSE: Forward Stepwise (0.81063)

Subset selection comparison interpretation:

The table above compares the three subset selection methods on the number of predictors chosen, BIC at the selected model size, and test prediction error. If all three methods agree on the same predictors, it suggests the selected variables have robust independent signal for burnout that is not sensitive to the direction of search. If the methods disagree, the differences are most likely driven by the correlated predictors identified in §4b, where the order of entry or removal changes which variables appear marginal. The method with the lowest test MSE is the most practically useful, but small differences in test MSE should not be over-interpreted given the CV variance seen in Chapter 6.

§5. Analysis

```
In [16]: # — Fit sklearn LinearRegression on full feature set —————
# sklearn is used for predictions and sklearn-based metrics.
# statsmodels OLS is used in parallel to get inference statistics
# (standard errors, t-stats, p-values, CIs) that sklearn does not provide.
```

```
mlr_model = LinearRegression(fit_intercept=True)
mlr_model.fit(X_train, y_train)

y_pred_train = mlr_model.predict(X_train)
y_pred_test  = mlr_model.predict(X_test)

train_r2     = round(r2_score(y_train, y_pred_train), 3)
test_r2      = round(r2_score(y_test, y_pred_test), 3)
train_rmse   = round(np.sqrt(mean_squared_error(y_train, y_pred_train)), 3)
test_rmse    = round(np.sqrt(mean_squared_error(y_test, y_pred_test)), 3)

print('Performance metrics (all to 3 decimal places):')
print(f' Train R-squared : {train_r2}')
print(f' Test  R-squared : {test_r2}')
print(f' Train RMSE      : {train_rmse}')
print(f' Test  RMSE      : {test_rmse}')
```

```
Performance metrics (all to 3 decimal places):
Train R-squared : 0.707
Test  R-squared : 0.707
Train RMSE      : 0.9
Test  RMSE      : 0.9
```

```
In [17]: # — Fit statsmodels OLS for full inference —————
# statsmodels gives us standard errors, t-statistics, p-values, and
# confidence intervals for each coefficient.
```

```
X_train_df = pd.DataFrame(X_train, columns=feature_names)
X_sm        = sm.add_constant(X_train_df)
ols_fit     = sm.OLS(y_train, X_sm).fit()

# — Coefficient table —————
```

```

ci          = ols_fit.conf_int(alpha=0.05)    # 95% confidence intervals

coef_table = pd.DataFrame({
    'Feature'      : feature_names,
    'Coefficient'   : [round(ols_fit.params[f], 3) for f in feature_names],
    'Std Error'     : [round(ols_fit.bse[f], 3) for f in feature_names],
    't-stat'        : [round(ols_fit.tvalues[f], 3) for f in feature_names],
    'p-value'       : [round(ols_fit.pvalues[f], 3) for f in feature_names],
    'CI Lower'      : [round(ci.loc[f, 0], 3) for f in feature_names],
    'CI Upper'      : [round(ci.loc[f, 1], 3) for f in feature_names],
})

print('Full coefficient table saved for Appendix C. Preview shown below:')
print(coef_table.to_string(index=False))

# Save coefficient table
coef_csv = os.path.join(RESULTS_DIR, 'coefficients_full_table_appendix.csv')
coef_table.to_csv(coef_csv, index=False)
print(f'\nSaved: {coef_csv}')

# Main chapter table: top coefficients by absolute value
coef_top_main = coef_table.sort_values('Coefficient', key=abs, ascending=False).head(25)
coef_top_csv = os.path.join(RESULTS_DIR, 'top_standardized_coefficients.csv')
coef_top_main.to_csv(coef_top_csv, index=False)
print()
print('Table 4.5: Top standardized coefficients by absolute magnitude.')
display(coef_top_main)
print(f'Saved: {coef_top_csv}')

```

Full coefficient table saved for Appendix C. Preview shown below:

Feature	Coefficient	Std Error	t-stat	p-value	CI Lower	CI Upper
age	-0.001	0.002	-0.273	0.785	-0.005	0.004
study_hours_per_day	0.007	0.004	1.586	0.113	-0.002	0.015
exam_pressure	0.231	0.004	53.295	0.000	0.222	0.239
academic_performance	0.001	0.003	0.258	0.796	-0.005	0.006
stress_level	0.583	0.005	108.043	0.000	0.572	0.593
anxiety_score	0.376	0.004	101.419	0.000	0.369	0.383
depression_score	0.310	0.003	89.277	0.000	0.303	0.317
sleep_hours	-0.338	0.003	-128.733	0.000	-0.343	-0.333
physical_activity	0.003	0.003	1.095	0.273	-0.002	0.008
social_support	-0.267	0.003	-98.207	0.000	-0.273	-0.262
screen_time	-0.017	0.005	-3.141	0.002	-0.027	-0.006
internet_usage	0.013	0.005	2.445	0.014	0.003	0.023
financial_stress	-0.003	0.003	-0.922	0.357	-0.009	0.003
family_expectation	-0.007	0.003	-2.371	0.018	-0.012	-0.001
gender_Male	-0.003	0.002	-1.405	0.160	-0.008	0.001
gender_Other	-0.001	0.002	-0.228	0.820	-0.005	0.004
academic_year_2	-0.000	0.003	-0.059	0.953	-0.006	0.006
academic_year_3	-0.001	0.003	-0.380	0.704	-0.007	0.005
academic_year_4	-0.001	0.003	-0.416	0.678	-0.007	0.005

Saved: ../results/ch04/coefficients_full_table_appendix.csv

Table 4.5: Top standardized coefficients by absolute magnitude.

	Feature	Coefficient	Std Error	t-stat	p-value	CI Lower	CI Upper
4	stress_level	0.583	0.005	108.043	0.000	0.572	0.593
5	anxiety_score	0.376	0.004	101.419	0.000	0.369	0.383
7	sleep_hours	-0.338	0.003	-128.733	0.000	-0.343	-0.333
6	depression_score	0.310	0.003	89.277	0.000	0.303	0.317
9	social_support	-0.267	0.003	-98.207	0.000	-0.273	-0.262
2	exam_pressure	0.231	0.004	53.295	0.000	0.222	0.239
10	screen_time	-0.017	0.005	-3.141	0.002	-0.027	-0.006
11	internet_usage	0.013	0.005	2.445	0.014	0.003	0.023

Saved: ../results/ch04/top_standardized_coefficients.csv


```
In [18]: # — Figure 4.3: Horizontal bar chart of standardized coefficients —————
# Coefficients from the scaled model are already on a comparable scale
# (standard deviations of X), so the bar lengths directly reflect relative
# importance. Positive bars mean higher burnout; negative bars mean lower burnout.

coef_sorted = coef_table.sort_values('Coefficient', key=abs, ascending=True)

colors = ['firebrick' if c > 0 else 'steelblue' for c in coef_sorted['Coefficient']]

fig, ax = plt.subplots(figsize=(9, 7))
ax.barh(coef_sorted['Feature'], coef_sorted['Coefficient'], color=colors)
ax.axvline(0, color='black', linewidth=0.8)
ax.set_xlabel('Coefficient (standardized units)', fontsize=11)
ax.set_title('Figure 4.3: Standardized Coefficients for All Predictors',
             fontsize=11, pad=10)
plt.tight_layout()
fig_path_43 = os.path.join(FIGURES_DIR, 'figure_4_3_coefficients.png')
plt.savefig(fig_path_43, dpi=150, bbox_inches='tight')
plt.show()
print(f'Saved: {fig_path_43}')
```

Figure 4.3: Standardized Coefficients for All Predictors



Saved: ../figures/ch04/figure_4_3_coefficients.png

§6. Diagnostics and Tuning

```
In [19]: # — 10-Fold Cross-Validation —————
# Cross-validation splits the training set into 10 folds, trains on 9 and
# evaluates on 1, cycling through all folds. This confirms that the train/test
# metrics are not driven by a lucky random split.

cv_scores = cross_val_score(
    estimator=LinearRegression(),
    X=X_train,
    y=y_train,
    cv=10,
    scoring='r2',
    n_jobs=-1
)

cv_mean = round(cv_scores.mean(), 3)
cv_std = round(cv_scores.std(), 3)

print('10-Fold CV R-squared scores (3 decimal places):')
for i, s in enumerate(cv_scores, 1):
    print(f'  Fold {i:2d}: {s:.3f}')
print()
```

```
print(f'CV Mean : {cv_mean}')
```

```
print(f'CV Std : {cv_std}')
```

10-Fold CV R-squared scores (3 decimal places):

```
Fold 1: 0.700
Fold 2: 0.709
Fold 3: 0.703
Fold 4: 0.714
Fold 5: 0.706
Fold 6: 0.710
Fold 7: 0.710
Fold 8: 0.707
Fold 9: 0.707
Fold 10: 0.708
```

CV Mean : 0.707

CV Std : 0.003

```
In [20]: # — Figure 4.4: Residuals vs Fitted Values —————
# Residuals = actual - predicted.
# A random scatter around the zero line suggests the linear model is appropriate.
# A fan shape suggests heteroscedasticity; a curve suggests a missed nonlinearity.
#
# A 3,000-point sample is used for display only. All residual statistics
# and diagnostics use the full training set.

residuals_train = y_train - y_pred_train


SAMPLE_N = 3000
rng = np.random.default_rng(RANDOM_STATE)
sample_idx = rng.choice(len(y_train), size=SAMPLE_N, replace=False)

fig, ax = plt.subplots(figsize=(8, 5))
ax.scatter(
    y_pred_train[sample_idx], residuals_train[sample_idx],
    alpha=0.25, s=12, color='steelblue'
)
ax.axhline(0, color='firebrick', linewidth=1.5, linestyle='--', label='Zero reference')
ax.set_xlabel('Fitted Values', fontsize=11)
ax.set_ylabel('Residuals', fontsize=11)
ax.set_title('Figure 4.4: Residuals vs Fitted Values', fontsize=11, pad=10)
ax.legend(fontsize=9)
plt.tight_layout()
fig_path_44 = os.path.join(FIGURES_DIR, 'figure_4_4_residuals.png')
plt.savefig(fig_path_44, dpi=150, bbox_inches='tight')
plt.show()
print(f'Saved: {fig_path_44}')
print(f'Note: figure shows {SAMPLE_N:,} sample points for readability.')
print('All residual statistics use the full training set.')

print('Interpretation note: the diagonal lower boundary appears because burnout score is 0')
print('This does not invalidate the model, but it shows that the 0 to 10 outcome scale is not linear')

```

Figure 4.4: Residuals vs Fitted Values



Saved: ../figures/ch04/figure_4_4_residuals.png

Note: figure shows 3,000 sample points for readability.

All residual statistics use the full training set.

Interpretation note: the diagonal lower boundary appears because burnout scores are bounded below at 0.

This does not invalidate the model, but it shows that the 0 to 10 outcome scale affects residual shape.

```
In [21]: # — Figure 4.5: Normal Q-Q Plot of Residuals —————
# Compares the distribution of residuals to a theoretical normal distribution.
# Points close to the diagonal line indicate approximately normal residuals.
# This figure uses all training residuals.


fig, ax = plt.subplots(figsize=(6, 6))

(osm, osr), (slope_qq, intercept_qq, _) = stats.probplot(
    residuals_train, dist='norm'
)

ax.scatter(osm, osr, alpha=0.3, s=6, color='steelblue', label='Residual quantiles')
x_ref = np.array([osm.min(), osm.max()])
ax.plot(x_ref, slope_qq * x_ref + intercept_qq,
        color='firebrick', linewidth=1.5, label='Normal reference line')

ax.set_xlabel('Theoretical Quantiles', fontsize=11)
ax.set_ylabel('Sample Quantiles (residuals)', fontsize=11)
ax.set_title('Figure 4.5: Normal Q-Q Plot of Residuals', fontsize=11, pad=10)
ax.legend(fontsize=9)
plt.tight_layout()
fig_path_45 = os.path.join(FIGURES_DIR, 'figure_4_5_qq.png')
plt.savefig(fig_path_45, dpi=150, bbox_inches='tight')
plt.show()
print(f'Saved: {fig_path_45}')
```

Figure 4.5: Normal Q-Q Plot of Residuals



Saved: ../figures/ch04/figure_4_5_qq.png

§7. Findings

```
In [22]: # — Table 4.6: Summary of MLR Results —————
# All values to exactly 3 decimal places as required by Chapter 4.

adj_r2 = round(ols_fit.rsquared_adj, 3)
f_stat = round(ols_fit.fvalue, 3)

# Top 5 predictors by absolute coefficient magnitude
top5 = coef_table.sort_values('Coefficient', key=abs, ascending=False).head(5)
top5_str = ', '.join(
    f"{row['Feature']} ({row['Coefficient']:.3f})"
    for _, row in top5.iterrows()
)

results_data = {
```

```

'Metric': [
    'Train R-squared',
    'Test R-squared',
    'Train RMSE',
    'Test RMSE',
    'Adjusted R-squared',
    'F-statistic',
    'CV R-squared (mean +/- std)',
    'Top 5 Predictors by |Coefficient|',
],
'Value': [
    f'{train_r2:.3f}',
    f'{test_r2:.3f}',
    f'{train_rmse:.3f}',
    f'{test_rmse:.3f}',
    f'{adj_r2:.3f}',
    f'{f_stat:.3f}',
    f'{cv_mean:.3f} +/- {cv_std:.3f}',
    top5_str,
]
}

results_df = pd.DataFrame(results_data)

print('Table 4.6: Multiple Linear Regression -- Key Results')
print('=' * 70)
print(results_df.to_string(index=False))
print('=' * 70)

```

Table 4.6: Multiple Linear Regression -- Key Results

```

=====
Metric
Value
Train R-squared
0.707
Test R-squared
0.707
Train RMSE
0.900
Test RMSE
0.900
Adjusted R-squared
0.707
F-statistic
17818.076
CV R-squared (mean +/- std)
0.707 +/- 0.003
Top 5 Predictors by |Coefficient| stress_level (0.583), anxiety_score (0.376), sleep
_hours (-0.338), depression_score (0.310), social_support (-0.267)
=====

```

```

In [23]: # — Save results to CSV —————
csv_path = os.path.join(RESULTS_DIR, 'mlr_results.csv')
results_df.to_csv(csv_path, index=False)
print(f'Table 4.1 saved to: {csv_path}')

```

```
mlr_key_results_csv = os.path.join(RESULTS_DIR, 'table_4_6_mlr_key_results.csv')
results_df.to_csv(mlr_key_results_csv, index=False)
print(f'Table 4.6 copy saved to: {mlr_key_results_csv}')
```

Table 4.1 saved to: ../results/ch04/mlr_results.csv

Table 4.6 copy saved to: ../results/ch04/table_4_6_mlr_key_results.csv

§8. Real-World Decision

The model shows that stress level, anxiety, and depression have the largest positive coefficients, while sleep hours and social support have the strongest negative coefficients. A university wellness office looking to act on these results should focus early outreach on students who score high on psychological distress indicators and low on protective factors such as sleep and support. Academic variables like exam pressure also contribute, so advising support during high-pressure periods would complement a mental-health-focused intervention. Gender and academic year show smaller effects, which means they are less useful on their own for targeting. Because the data source does not clearly confirm how the records were collected or created, the specific coefficient values should not be treated as direct evidence for real student populations without outside validation.

§9. Caveats and Limitations

Even after flagging high-VIF features, some multicollinearity remains between stress, anxiety, and depression, which are conceptually overlapping constructs. This means the individual coefficients for those variables should be interpreted with caution because part of their effect is shared. Since the data source does not clearly confirm how the records were collected or created, the strong R-squared values should be treated as evidence of model performance on this dataset, not as confirmed real-world student behavior. As with Chapter 3, the model captures associations, not causal relationships, so higher stress scores being linked to higher burnout does not prove that reducing stress alone would directly lower burnout.

```
In [24]: # — Final summary —————

print('Chapter 4 complete. Files written:')
print()

print('Main chapter figures')
for p in [fig_path_41, fig_path_42, fig_path_43, fig_path_44, fig_path_45]:
    print(f' {p}')

print()

print('Appendix C figures')
```

```

for p in [fig_app_c41, fig_app_c42, fig_app_c43]:
    print(f' {p}')

print()

print('Main results and tables')
for p in [
    csv_path,
    mlr_key_results_csv,
    vif_high_csv,
    comparison_csv,
    coef_top_csv
]:
    print(f' {p}')

print()

print('Appendix C tables')
for p in [
    vif_full_csv,
    subset_csv,
    best_subset_csv,
    bwd_csv,
    coef_csv
]:
    print(f' {p}')

```

Chapter 4 complete. Files written:

Main chapter figures

```

../figures/ch04/figure_4_1_vif.png
../figures/ch04/figure_4_2_aic_bic_selection.png
../figures/ch04/figure_4_3_coefficients.png
../figures/ch04/figure_4_4_residuals.png
../figures/ch04/figure_4_5_qq.png

```

Appendix C figures

```

../figures/ch04/appendix_c_4_1_greedy_best_subset_bic.png
../figures/ch04/appendix_c_4_2_backward_stepwise_bic.png
../figures/ch04/appendix_c_4_3_forward_stepwise_bic.png

```

Main results and tables

```

../results/ch04/mlr_results.csv
../results/ch04/table_4_6_mlr_key_results.csv
../results/ch04/vif_high_predictors.csv
../results/ch04/table_4_4_subset_comparison.csv
../results/ch04/top_standardized_coefficients.csv

```

Appendix C tables

```

../results/ch04/vif_full_table_appendix.csv
../results/ch04/subset_selection.csv
../results/ch04/greedy_best_subset_table_appendix.csv
../results/ch04/backward_stepwise_table_appendix.csv
../results/ch04/coefficients_full_table_appendix.csv

```